

BigQuery STRUCT

TABLE OF CONTENTS

1. Overview
2. Nested and Repeated Schema
3. Effect of Use
4. Usability
5. Conclusion

Overview

Objective

In an environment where BigQuery is used as a DB service,
proposing cost-saving measures
through the columnar storage features of BigQuery and its supported schemas.

BigQuery

Cloud service for storing and analyzing petabyte-scale data,
provided to external users by Google

Based on [Dremel](#),
it features **Columnar Storage** and SQL tree architecture

BigQuery

Nested and Repeated Schema

BigQuery Nested and Repeated Schema

Denormalization

When BigQuery is denormalized, performance is optimized
Implementing Nested and Repeated to achieve denormalization

Advantages

- Performance Improvement
- Efficient data storage capacity
- Flexible adaptation to schema changes

Columnar Storage

Storing tables at the denormalized column level to achieve storage efficiency

BigQuery Nested and Repeated Schema

Nested

BigQuery : RECORD

Container that have Every **Type(Required)** and field_name(Optional) sorted
STRUCT data structure

Similar to Dictionary(Python), **but with fixed Key(Schema exists)**

→ **You can load data into a single column by bundling the data into one data struct**

Possible to nest up to depth of 15 levels

→ Same storage size does not matter of **Nest(Storage based Pricing)**

| 상품명 | 가격 | 수량 |
|-----|-----|----|
| A | 100 | 1 |
| C | 300 | 2 |
| A | 100 | 1 |
| B | 200 | 2 |
| C | 300 | 3 |

of Column : 3

| 주문 |
|--------------------------|
| 상품명: 'A', 가격: 100, 수량: 1 |
| 상품명: 'C', 가격: 300, 수량: 2 |
| 상품명: 'A', 가격: 100, 수량: 1 |
| 상품명: 'B', 가격: 200, 수량: 2 |
| 상품명: 'C', 가격: 300, 수량: 3 |

of Column : 1
Actual Form

| 주문 | | |
|--------|-------|-------|
| 주문.상품명 | 주문.가격 | 주문.수량 |
| A | 100 | 1 |
| C | 300 | 2 |
| A | 100 | 1 |
| B | 200 | 2 |
| C | 300 | 3 |

of Column : 1
BigQuery Visualization

BigQuery Nested and Repeated Schema

Repeated

BigQuery : REPEATED

Array with same data type

Similar to Array(Python)

If ARRAY has any NULL value, it will Return NULL

→ IGNORE NULLS to avoid

BigQuery Visualization of Array is as follows

| 성 | 이름 | 연락처 | 나이 | 상품명 |
|---|----|----------|----|-----|
| 홍 | 은지 | 010-3333 | 30 | A |
| 홍 | 은지 | 010-3333 | 30 | C |
| 이 | 유나 | 010-4444 | 29 | A |
| 이 | 유나 | 010-4444 | 29 | B |
| 이 | 유나 | 010-4444 | 29 | C |

Original

| 성 | 이름 | 연락처 | 나이 | 상품명 |
|---|----|----------|----|-----------|
| 홍 | 은지 | 010-3333 | 30 | [A, C] |
| 이 | 유나 | 010-4444 | 29 | [A, B, C] |

Actual Form

| 성 | 이름 | 연락처 | 나이 | 상품명 |
|---|----|----------|----|-----|
| 홍 | 은지 | 010-3333 | 30 | A |
| | | | | C |
| 이 | 유나 | 010-4444 | 29 | A |
| | | | | B |
| | | | | C |

BigQuery Visualization

BigQuery Nested and Repeated Schema

Nested and Repeated

RECORD(STRUCT) + ARRAY

Denormalization using STRUCT with ARRAY

Apply array then use struct method

Original

| 성 | 이름 | 연락처 | 나이 | 상품명 | 가격 | 수량 |
|---|----|----------|----|-----|-----|----|
| 홍 | 은지 | 010-3333 | 30 | A | 100 | 1 |
| 홍 | 은지 | 010-3333 | 30 | C | 300 | 2 |
| 이 | 유나 | 010-4444 | 29 | A | 100 | 1 |
| 이 | 유나 | 010-4444 | 29 | B | 200 | 2 |
| 이 | 유나 | 010-4444 | 29 | C | 300 | 3 |

Nested



| 성 | 이름 | 연락처 | 나이 | 주문 | | |
|---|----|----------|----|--------|-------|-------|
| | | | | 주문.상품명 | 주문.가격 | 주문.수량 |
| 홍 | 은지 | 010-3333 | 30 | A | 100 | 1 |
| 홍 | 은지 | 010-3333 | 30 | C | 300 | 2 |
| 이 | 유나 | 010-4444 | 29 | A | 100 | 1 |
| 이 | 유나 | 010-4444 | 29 | B | 200 | 2 |
| 이 | 유나 | 010-4444 | 29 | C | 300 | 3 |

Nested and Repeated

| 성 | 이름 | 연락처 | 나이 | 주문 | | |
|---|----|----------|----|--------|-------|-------|
| | | | | 주문.상품명 | 주문.가격 | 주문.수량 |
| 홍 | 은지 | 010-3333 | 30 | A | 100 | 1 |
| | | | | C | 300 | 2 |
| 이 | 유나 | 010-4444 | 29 | A | 100 | 1 |
| | | | | B | 200 | 2 |
| | | | | C | 300 | 3 |



- Efficient Data Storage
- Form of GROUP BY

Effect of Use

Effect of Use

Test Data

DB data accumulated over 3 days

Duration :

Comparison

Applying a new schema to a total of 80 tables

22 Tables stay as current

37 Tables applying single STRUCT to achieve capacity efficiency

21 Tables applying several STRUCT to achieve capacity efficiency and table unification

| Unification | 4 → 1 | 3 → 1 | 2 → 1 |
|-------------------|--------|-------|-------|
| # of result Table | 12 → 3 | 3 → 1 | 6 → 3 |

Total 66 Tables are created for capacity comparison

Effect of Use

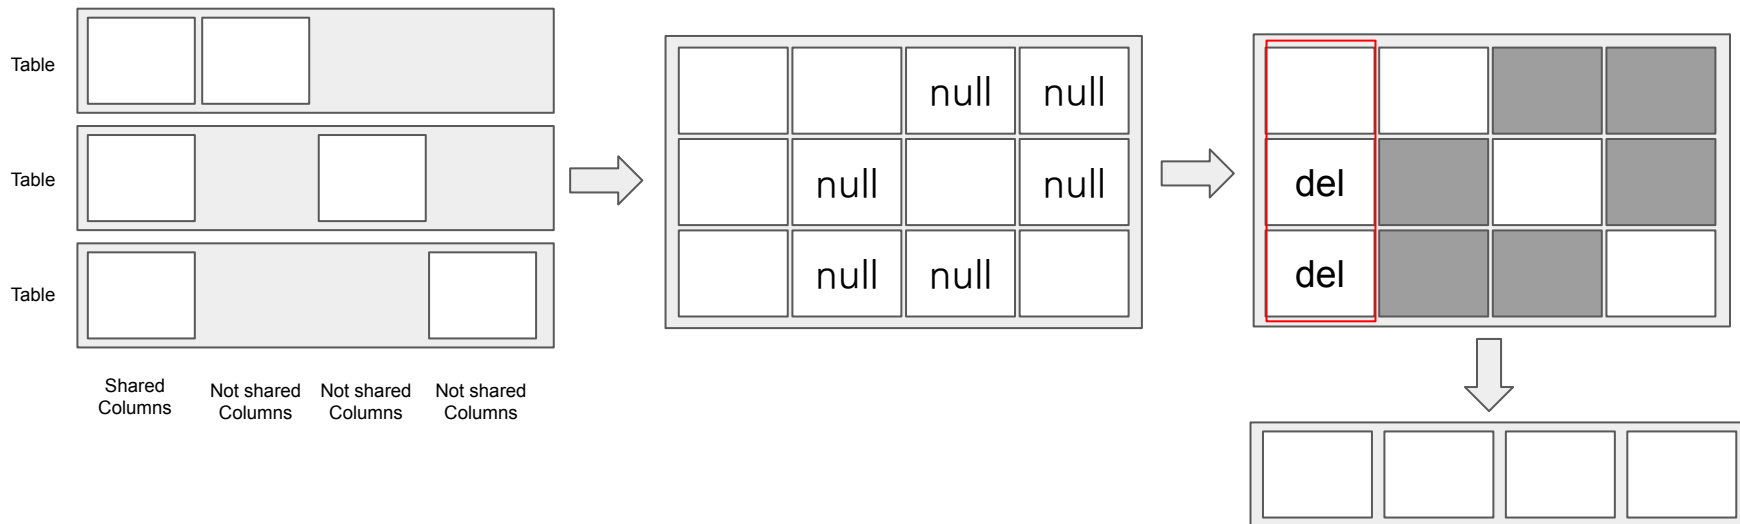
Table Unification using STRUCT

Table Unification based on Shared Columns

Make STRUCT with Not Shared Columns

Null Value does not take up space, because of columnar storage

Storage efficiency and table unification both can be achieved



Effect of Use

결과

| Whole Tables | As Is | To Be | Reduction Amount |
|--------------|--------------|------------------------|------------------|
| # of Tables | 80 | 66 | 14 |
| Capacity | TB (100%) | TB (48.02%) | TB (51.98%) |
| Price | ₩ | ₩ | ₩ |

| W/O same Tables | As Is | To Be | Reduction Amount |
|-----------------|--------------|------------------------|------------------|
| # of Tables | 58 | 44 | 14 |
| Capacity | TB (100%) | TB (43.79%) | TB (56.21%) |
| Price | ₩ | ₩ | ₩ |

Effect of Use


Result

Average storage size decline rate of denormalized table : 47.66 %

Comparison between denormalization and unification table

Comparison between 3 source table and results

| Table | Source Table | Denormalization | Unification(denormalized and unified) |
|---------|-------------------|--|--|
| Table 1 | MB (100%) | MB (40.62%) | - |
| Table 2 | MB (100%) | MB (36.06%) | |
| Table 3 | GB (100%) | MB (35.26%) | |
| Total | GB (100%) - | GB Compared to Original (37.21%) - | GB Compared to Original (27.27%) Compared to Denormalization(73.28%) |



※ 한 record에 여러 Not Shared Column들이 조합되면 압축 효율이 높아질 수 있음.

Usability

Usability

UNNEST

need to use UNNEST function when using RECORD(STRUCT) + ARRAY

※ Access is possible in the form of “Nested_field.*” to reference all fields within the nested structure

#1


```
SELECT *  
FROM `project.dataset.table`  
CROSS JOIN  
    UNNEST(nested_field)
```

#2

```
SELECT * EXCEPT(nested_field)  
FROM `project.dataset.table`, UNNEST(nested_field)
```

#3 access field with reference

```
SELECT ALIAS.target_nested_field  
FROM `project.dataset.table`, UNNEST(nested_field) AS ALIAS
```



※ UNNEST 할 때 nested_field를 제외하지 않고 전체 SELECT 시 위 그림과 같이 SELECT 하는 경우가 있어 주의할 것

Usability

Criteria for RECORD Configuration

1. Exclusion of Partition Columns
The leaf fields of a STRUCT cannot be used as partitions.
Assumes partition usage for scan speed improvement
2. Conversion of Primary Key and Measure into STRUCT
Measure fields are transformed into RECORD for non-code-based keys, considering the low likelihood of duplication.
3. Grouping Similar Types
Aim to improve compression ratios through encoding by grouping similar types
4. Frequently Used Columns are placed outside the STRUCT for usability

Usability

Comparison by possible combinations of external columns from STRUCT

Should consider 1) Too many cases and 2) impossible cases for nothing

Comparison by DISTINCT combinations of external columns from STRUCT

The compression ratio seems to follow a linear relationship with the number of distinct combinations of external columns from STRUCT.

However, there appears to be a phenomenon where the graph inverts after a certain level of compression.

| # of external columns | possible combinations | DISTINCT combinations | record reduction rate | storage size | storage reduction rate |
|-----------------------|---------------------------|-----------------------|-----------------------|--------------|------------------------|
| 0 (Source) | X | 408106 | X | 55.05MB | X |
| 1 | 138,000,000,000 | 408106 | 0% | 55.05MB | 0% |
| 2 | 1,270,000,000,000,000 | 357060 | 13% | 50.03MB | 9% |
| 3 | 1,010,000,000,000,000,000 | 247168 | 39% | 35.70MB | 35% |
| 4 | 13,158 | 6247 | 98% | 28.45MB | 48% |
| 5 | 1 | 53 | 100% | 33.83MB | 39% |

Usability

Optimization on Single STRUCT

$\min[\text{possible combinations of external columns from STRUCT} * \# \text{ of external columns} + \# \text{ of every Records} * \text{internal columns}]$

→ Minimize the white area of graph


$\approx \max[(\# \text{ of every Records} - \text{possible combinations of external columns from STRUCT}) * \# \text{ of external columns}]$

→ Maximize the grey area of graph

→ Instead of possible combinations of external columns of STRUCT,
DISTINCT combinations of external columns of STRUCT
is expected to be more applicable for use

Limitation

- Combination
 - Dependency between variables are not considered
 - using DISTINCT case for real life applicatino
- Storage size : data type size need to be considered
eg) text type with long length
- single STRUCT denormalization is only considered



Usability

Optimization in mathematical form

N = # of columns

X = external columns of STRUCT $\rightarrow N - X$ = internal columns of STRUCT

n = # of every records

A = # of possible combinations of external column from STRUCT

a = # of DISTINCT combinations of external column from STRUCT

Condition : $a \rightarrow A$ as $n \rightarrow \infty$

$$\min[A*X + n*(N-X)]$$

$$= \min[A*X + n*N - n*X] \approx \min[A*X - n*X]$$

$\because n*N : \text{Constant}$


$$= \min[(A - n)*X] = -\max[(A - n)*X]$$

$$= \max[(n - A)*X]$$

$$\therefore \min[A*X + n*(N-X)] \approx \max[(n - A)*X]$$

\because by Condition : $a \rightarrow A$

$$\therefore \min[a*X + n*(N-X)] \approx \max[(n - a)*X]$$



Conclusion

Conclusion

Summary

- Depending on the table, the suitable form of nested columns may vary. Therefore, it is important to note that the criteria may differ for each table.
- Recommended to use a single STRUCT per table. If multiple STRUCTs are used, it is advised to configure them in a way that the STRUCTs are not interlinked
- For partitions and frequently used columns, placing them in external columns is recommended. Measures are recommended to be placed in internal columns of the STRUCT
- When structuring, it is recommended to consider the number of distinct combinations of external columns in a single STRUCT as a guideline for usage

Appendix

BigQuery Pricing

BigQuery Pricing

Storage based Pricing

Cost of storing data loaded into BigQuery

Active Storage

Every Table or Table Partitions modified in the last 90 days

Long-Term Storage

Every Table or Table Partitions that have not been modified in the last 90 days

50% price reduction automatically

No difference in performance, durability, availability between active and long-term

| Storage | US Region | Tokyo Region | Details |
|-----------|----------------|----------------|-----------------------|
| Active | \$0.020 per GB | \$0.023 per GB | Free 10GB every month |
| Long-Term | \$0.010 per GB | \$0.016 per GB | Free 10GB every month |

BigQuery Pricing

Analysis Pricing

cost to process queries, including SQL queries, user-defined functions, scripts, and certain data manipulation language (DML) and data definition language (DDL) statements

On-demand Pricing

charged for the number of bytes processed by each query
first 1 TiB of query data processed per month is free

Flat-rate Pricing

Purchasing Virtual CPU 'SLOT'

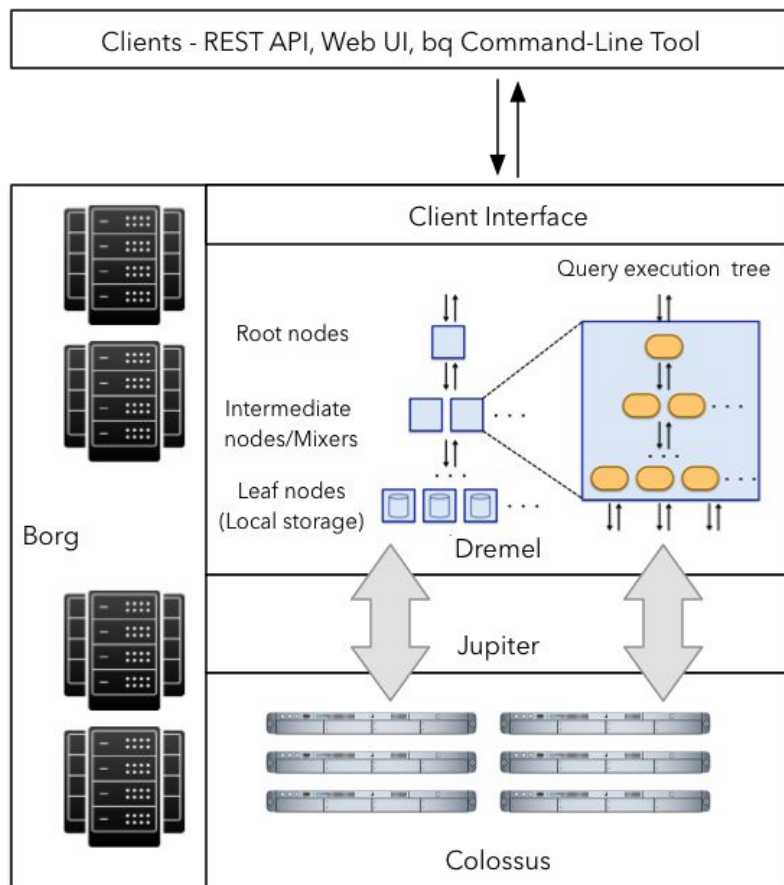
Variable slots: Committed for the first 60 seconds.

Monthly: Committed for the first 30 days.

Annual: Committed for 365 days.

BigQuery Columnar Storage

BigQuery Columnar Storage



Borg

large-scale cluster management system
allocates server resources

Dremel

Columnar Storage
tree-based execution structure

Jupiter

Network Infrastructure
Bisectional Bandwidth

- 1 peta bit per second

- Decreased importance of locality (region)
- Reduced significance of lag considerations.

Colossus

Successor to the Google File System (GFS)
Capacitor : Columnar Storage Format

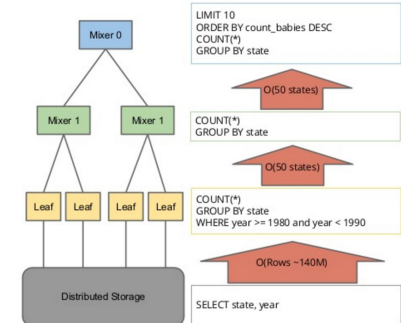
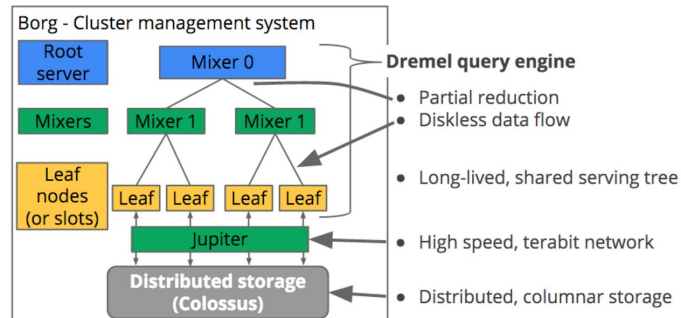
BigQuery Columnar Storage

Dremel Engine


- database with the characteristic of storing data at the column level
- combining columnar storage and a tree architecture for SQL queries, achieving exceptionally fast speeds
- efficiency of data scanning is high due to the columnar storage method
- advantageous for read-centric OLAP tasks rather than OLTP

Tree Architecture Distribution(Execution Tree)


- Root Server(Mixer 0), Intermediate Server(Mixer1).....Leaf Server(Leaf)
- Distributes SQL queries in a tree structure to distributed machines and achieves fast processing.
- It processes data directly in RAM, without using disks except for leaf nodes, resulting in rapid speed.



BigQuery Columnar Storage



record-oriented



column-oriented

Record Oriented
FILE 1 : 1;Kim;F;32,
FILE 2 : 2;Nam;F;31,
FILE 3 : 3;Hong;M;30

Column Oriented
FILE 1 : 1:Kim, 2:Nam, 3:Hong
FILE 2 : 1:F, 2:F, 3:M
FILE 3 : 1:32, 2:31, 3:30

BigQuery Columnar Storage

Properties

- Traffic Minimization
Cost savings occur because only the necessary columns are queried.
- High Compression Rate
Gathering and storing similar data types makes compression easier.
Run Length Encoding(RLE), Dictionary Encoding

Encoding

Repetition and Definition Level

Concepts used when reconstructing a column storage record.

Combine by calling each columns→ Efficient query capacity achieved

Stored together for each Block fragment (by column)

Capacitor

BigQuery's Column-Based Storage Format

Collects various statistical information and stores it in Colossus.

Estimates and generates optimization models based on the gathered information.

BigQuery Columnar Storage

Other Characteristics

- No Key, No Index(Full Scan Only)

Avoid full scan with Partitioning

- Eventual Consistency

Replicating data to three data centers

Typically, it is reflected immediately, but in some situations, it may take a few minutes.

